

Why Verbs Are Harder to Learn Than Nouns: A Quantitative Approach

Anonymous CogSci submission

Abstract

Include no author information in the initial submission, to facilitate blind review. The abstract should be one paragraph, indented 1/8 inch on both sides, in 9 point font with single spacing. The heading 'Abstract' should be 10 point, bold, centered, with one line of space below it. This one-paragraph abstract section is required only for standard six page proceedings papers. Following the abstract should be a blank line, followed by the header 'Keywords' and a list of descriptive keywords separated by semicolons, all in 9 point font, as shown below.

Keywords: Add your choice of indexing terms or keywords; kindly use a semi-colon; between each term.

By the time they are two years old, children fill out their vocabularies with hundreds of words - many of these are nouns like "dog" and "ball", but others are social performatives or proto-verbs like "bye" and "up" (Braginsky, Yurovsky, Marchman, & Frank, 2019; Tardiff et al., 2008).

(Behavior experiment evidence suggests that children learn a disproportionate amount of nouns in their early age.) The core difference is that verbs are more variable than nouns in various ways (Gentner). First, verbs are more variable in their semantics. One verb can be used to mention several actions, and one action can always be described by more than one verb. Second, verb meanings are more variable across languages, while nouns are relatively stable cross-linguistically. Finally, verbs are more variable in terms of word forms. More cognitive capacity is needed for children in order to figure out that different verb forms actually correspond to the same concept of verb.

To explain the phenomena from theory level, we examine different learning theories.

One possibility is that children learn cross-modal mappings directly from the co-occurrence statistics between the words they hear and their observations of the world around them. Although the referent of a word may be ambiguous in any individual context, children could learn a word's meaning by identifying what is consistent across the many contexts of its use (Siskind, 1996; Yu & Smith, 2007). That is, a child could learn that "ball" generally refers to round toys because it is frequently used when they are around. Evidence from a number of labs now shows that children and even infants can use cross-situational statistical information to learn the meanings of concrete nouns in simplified contexts (Smith & Yu, 2008; Vlach & Johnson, 2013; Suanda et al., 2014).

However, there are reasons to think that children do not learn all of these words by tracking cross-modal occurrence information. First, referential ambiguity in the world may be significantly more complex than the ambiguity faced by children in these lab studies (Medina et al., 2013; although c.f. Yurovsky et al., 2012). One extreme case of this is that words may refer not to co-present objects, but instead to things in the past or future. As Gleitman (1990) points out, a child's caregiver is unlikely to say "I am opening the door!" as they come home from work, but instead something like "It's freezing outside!". Second, studies showing successful cross-situational learning of non-noun meanings are rare, and successes have been found only in older children with significant extra scaffolding (Childers & Paik, 2009; Scott & Fisher, 2012).

An alternative possibility is that children learn meanings not from cross modal correspondence but from the structure of information within individual modalities. An intriguing body of evidence consistent with this possibility comes from the semantics of children who do not have access to visual information. Blind children—who cannot learn the meanings of words by mapping them onto their visual referents—have very similar semantics for highly visual words as sighted children (Landau, et al., 1985). Even sighted children develop an early understanding of the category of color words before they learn the meanings of the words themselves. That is, two-year-olds know that the right answer for a question like "what color is this" is "blue" or "green" rather than "dog" even if they cannot correctly identify which color is the right answer (Bartlett, 1978). Statistical models trained on corpora of English language are able to recover a striking amount of information about the meanings of words, even without the complex machinery that has enabled modern machine learning models to be so powerful (Lund & Burgess, 1996; Landauer & Dumais, 1997). But how could children ground these linguistic "meanings" in the physical world; how could these uni-modal meanings be used for cross-modal mapping without tracking cross-modal co-occurrence?

Alignment Theory

Recently, Roads and Love (2020) proposed a novel and interesting learning theory, which ingeniously excludes any use of cross-modality co-occurrence statistics - they view word learning as the unsupervised alignment of multiple concep-

tual systems (e.g., visual system, language system, acoustic system). [An example: dog, cat, tiger]. The key insight is that different sources of input should produce similar conceptual systems because sources are different viewpoints of the same underlying reality. If structural idiosyncrasies present in one system are qualitatively mirrored in the other system, then it is possible to align the two systems. Technically, the alignment is done in two steps. At step one, the learner infers the distribution of concepts within each independent modality by measuring the similarity of each pair of concepts. Once the learner has a good sense of the similarity structures of concepts within each conceptual system, at step two, the mapping of all concepts across systems can be set up at one go: Roads and Love discovered that the alignment correlation of two systems positively correlates with the mapping accuracy, which implies that the alignment correlation can be used as a hint to find the perfect mapping [one figure here]. [brief discussion: 1. Both cross-situational learning and the alignment idea are unsupervised methods. 2. The core difference is that in alignment theory, no cross-modality co-occurrence statistics is needed.] We use this alignment idea as an alternate framework to explore the verb problems from before

For more information on citations in RMarkdown, see [here](#).

Study 0

As we mentioned before, one of Roads and Love's most important findings is that the alignment correlations of systems positively correlate with the mapping accuracy of concepts, which suggests that better mapping usually co-occurs with higher alignment correlation.

We first verify this finding by replicating Roads and Love's experiment of noun concept alignment on a new dataset, using new embedding models. We demonstrated the alignment across the visual system and the language system.

Footnotes

Indicate footnotes with a number¹ in the text. Place the footnotes in 9 point type at the bottom of the page on which they appear. Precede the footnote with a horizontal rule.² You can also use markdown formatting to include footnotes using this syntax.³

Figures

All artwork must be very dark for purposes of reproduction and should not be hand drawn. Number figures sequentially, placing the figure number and caption, in 10 point, after the figure with one line space above the caption and one line space below it. If necessary, leave extra white space at the bottom of the page to avoid splitting the figure and figure caption. You may float figures to the top or bottom of a column, or set wide figures across both columns.

¹Sample of the first footnote.

²Sample of the second footnote.

³Sample of a markdown footnote.

Two-column images

You can read local images using png package for example and plot it like a regular plot using grid.raster from the grid package. With this method you have full control of the size of your image. **Note: Image must be in .png file format for the readPNG function to work.**

You might want to display a wide figure across both columns. To do this, you change the `fig.env` chunk option to `figure*`. To align the image in the center of the page, set `fig.align` option to `center`. To format the width of your caption text, you set the `num.cols.cap` option to 2.

One-column images

Single column is the default option, but if you want set it explicitly, set `fig.env` to `figure`. Notice that the `num.cols` option for the caption width is set to 1.



Figure 2: One column image.

R Plots

You can use R chunks directly to plot graphs. And you can use latex floats in the `fig.pos` chunk option to have more control over the location of your plot on the page. For more information on latex placement specifiers see [here](#)




Figure 3: R plot

Tables

Number tables consecutively; place the table number and title (in 10 point) above the table with one line space above the caption and one line space below it, as in Table 1. You may float tables to the top or bottom of a column, set wide tables across both columns.



Figure 1: This image spans both columns. And the caption text is limited to 0.8 of the width of the document.

You can use the xtable function in the xtable package.

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.03	0.10	0.3	0.74
x	2.02	0.10	21.1	0.00

Table 1: This table prints across one column.

Acknowledgements

Place acknowledgments (including funding information) in a section at the end of the paper.

References

- Braginsky, M., Yurovsky, D., Marchman, V. A., & Frank, M. C. (2019). Consistency and variability in children’s word learning across languages. *Open Mind*, 3, 52–67.
- Gleitman, L. (1990). The structural sources of verb meanings. *Language Acquisition*, 1(1), 3–55.
- Siskind, J. M. (1996). A computational study of cross-situational techniques for learning word-to-meaning mappings. *Cognition*, 61(1-2), 39–91.
- Tardiff, T., Fletcher, P., Liang, W., Zhang, Z., Kaciroti, N., & Marchman, V. (2008). Baby’s first ten words. *Development Psychology*, 44(4), 929–938.
- Yu, C., & Smith, L. B. (2007). Rapid word learning under uncertainty via cross-situational statistics. *Psychological Science*, 18(5), 414–420.